

Gary Miller

United States | Remote /Open to Relocation
[LinkedIn](#) | [GitHub](#) | [GM Cubed Designs](#)

Professional Summary

Founder & Full-Stack / AI Engineer and US Navy Veteran with 7+ years of experience architecting full-stack web applications, multi-agent AI systems, and local LLM orchestration pipelines. Skilled in RAG architectures, deterministic data processing, model quantization, and AI red teaming/RLHF evaluation. Proven track record leading client-facing engineering efforts, managing stakeholder expectations, and building privacy-focused, enterprise-grade AI applications. Brings rigorous Navy operational discipline and process control to fast-paced development environments.

Core Competencies & Technical Skills

- **AI & Machine Learning:** Local LLM Orchestration (Ollama, Hugging Face), Multi-Agent Frameworks (CrewAI), RAG Architectures, Vector Databases (Qdrant), Model Quantization, AI Red Teaming & RLHF, Prompt Engineering.
- **Languages & Core Tech:** Python, TypeScript, JavaScript, C++, C, SQL, Ruby.
- **Backend & Infrastructure:** FastAPI, Node.js, React, Next.js, PostgreSQL (TimescaleDB), Redis, Ruby on Rails, Docker, AWS.
- **Methodologies & Safety:** System Architecture, Data Pipelines, Agile Development, Continuous Integration / Deployment (CI/CD), Human-in-the-Loop (HITL) Evaluation.

Professional Experience

G M Cubed Designs | Founder & Full-Stack / AI Engineer | Apr 2019 – Present

- **Engineered full-stack solutions and AI web tools** that generated a **56.27% YoY revenue growth** by scaling client digital pipelines, search optimization, and conversion rates.
- **Collaborated directly with client teams and stakeholders** to translate technical specifications into scalable backend services, custom web architectures, and automated data workflows.
- **Architected local LLM orchestration frameworks and RAG setups** utilizing Ollama and Qdrant to enable secure, privacy-compliant document retrieval for client applications.
- **Managed full lifecycle delivery** across multiple concurrently deployed web platforms, applying rigorous testing and performance monitoring to maintain high system uptime.
- **Accelerated software delivery velocity** using modern AI-assisted workflows (Cursor, Claude Code) to build, test, and ship full-stack web applications.

US Navy | Aviation Structural Mechanic | Feb 2007 – Jun 2010

- Maintained critical structural, hydraulic, and flight-control systems on naval aircraft under strict operational deadlines, executing precision repairs and quality assurance with zero safety compromises.

Select Technical Projects

- **[Sidekick LLM:](#)**
Desktop Knowledge Management System | [GitHub](#)
 - **Architected a local desktop application** using Electron, React, and SQLite, enabling local LLM query execution over private vault files.
 - **Integrated federated retrieval and local source-grounding** to eliminate third-party API exposure and ensure offline data privacy.
 - **Validation Signal:** Deployed locally across active client desktop environments for offline document processing.
- **[SYNAPSE:](#)**
Multi-Agent AI Orchestration & Asynchronous Data Pipeline | [Github](#)
 - **Architected a decoupled multi-agent orchestration framework** using Python, FastAPI, and Redis to handle real-time context retrieval and persistent agent memory.
 - **Integrated Qdrant vector search and TimescaleDB time-series storage** to process streaming inputs across autonomous agents with low-latency execution.
 - **Validation Signal:** Built as an open-source framework and production-ready local pipeline setup for high-throughput AI workloads.
- **[Investment Matrix:](#)**
Quantitative Financial Modeling & Crypto Backtesting | [GitHub](#)
 - **Architected an institutional-grade financial modeling platform** using Python, Flask, and PostgreSQL to execute real-time market data ingestion and technical analysis pipelines.
 - **Engineered custom backtesting engines and risk metrics** to analyze crypto price action and historical volatility across dynamic streaming data feeds.
 - **Validation Signal:** Open-source quantitative framework supporting automated indicators and time-series feature engineering for algorithmic trading models.
- **[Holdem-Trainer:](#)**
High-Fidelity Logic Engine & Game Theory Simulation | [GitHub](#)
 - **Architected a Texas Hold'em training framework** with a high-fidelity validation suite featuring 700+ custom unit tests to proactively detect regressions and verify hand evaluations.
 - **Engineered deterministic hand evaluation logic and poker probability models** using Python to calculate exact pot odds and EV decisions under strict execution limits.
 - **Validation Signal:** Fully tested, production-ready simulation harness built with full test coverage across complex multi-way hand distributions.

Technical Writing & Publications

Author / Contributor as Atlas Anomalous | [Medium](#)

- Authored "Quantization in AI: Mastering Precision for Enterprise Efficiency," detailing advanced optimization architectures for deploying massive LLMs at scale, including INT8/INT4 precision scaling, Post-Training Quantization (PTQ), and Quantization-Aware Training (QAT) for edge AI environments.
- Authored comprehensive technical guides on full-stack development, including deep dives into Big O Notation and Asymptotic Analysis, Metaprogramming in Ruby, and API Integration.

Education

- **Harvard CS50** | Computer Science Online Course (2023)
- **App Academy** | Bootcamp in Full Stack Software Engineering (2019)
- **Austin Community College** | Associate Degree Economics (2013 – 2015)

Certifications

- **Lean Six Sigma Green Belt (SSGB)** | American Society of Quality (ASQ) 2015
- **Introduction to Cloud 101** | Amazon Web Services (AWS) 2024